



Mining Proportional Patterns in Online Retail Data

Your Name
VRXXXXXX

Second Name
VRXXXXXX

1 Introduction

Frequent Itemset Mining (FIM) has long been a cornerstone of market basket analysis, providing retailers with the ability to discover hidden associations between products purchased together. The evolution of these techniques led to the development of the FP-Growth algorithm, which revolutionized the field by adopting a "divide-and-conquer" approach that avoids the prohibitively expensive candidate generation typical of earlier methods like Apriori. However, in the era of massive online retail datasets, two primary challenges emerge: computational scalability and the inherent limitation of binary data.

1.1 The Challenge of Scale and the Long Tail

Modern retail databases often reach sizes where memory and computational costs on a single machine become unsustainable. Furthermore, web-based retail data frequently exhibits a "long-tail" distribution, where a massive number of items have low individual support but represent a significant portion of overall revenue when considered collectively. To address this, the Parallel FP-Growth (PFP) framework was introduced, enabling the distribution of mining tasks across a cluster of machines using a five-step MapReduce-based pipeline. PFP achieves virtually linear speedup by sharding the transaction database into independent tasks, eliminating the need for inter-machine communication during the mining process.

1.2 Moving Beyond Binary Mining

Despite the efficiency of PFP, traditional FIM frameworks are largely limited to the binary presence or absence of items, effectively ignoring the purchased quantities. In a retail context, understanding the exact proportions of items bought together (e.g., a "2-for-4" ratio of related products) provides deeper insights for bundle recommendations, targeted upselling, and inventory management.

1.3 The Proposed Extension: Proportional Analysis

This project proposes a structural extension to the PFP framework designed to extract proportional quantitative relationships. By implementing a custom "Bijective Token Indexing" phase, we adapt the PFP distributed engine to handle quantitative tuples without modifying its core prefix-tree logic. The primary contribution of this work lies in the post-mining analytical layer: we apply a Greatest Common Divisor (GCD) decomposition to identify the Canonical Shape of frequent patterns—the irreducible proportional ratio of items. This logic allows us to bridge the gap between tree-based mining and structural graph analysis, culminating in a Proportional Relationship Graph. By applying transitive reduction to these results, we eliminate redundant informational pathways, providing an optimized, hierarchical atlas of consumer purchase trajectories that highlights the structural evolution of consumer behavior.

2 Formal Methodology: Quantitative Itemsets, Canonical Shapes, and Scaling Factors

To rigorously analyze co-occurrences in retail data while accounting for purchased quantities, we establish a formal mathematical framework [1, 2]. Let $I = \{i_1, i_2, \dots, i_m\}$ be the universe of distinct items [3]. A **quantitative itemset** X is defined as a set of ordered pairs representing items and their specific purchased quantities: $X = \{(i_j, q_j) \mid i_j \in I, q_j \in \mathbb{Z}^+\}$ [3].

For any frequent quantitative pattern $P \in \mathcal{P}$ characterized by the quantity set $\{q_1, q_2, \dots, q_k\}$, we define the **Scaling Factor** α as the greatest common divisor of those quantities: $\alpha = \gcd(q_1, q_2, \dots, q_k)$ [4]. This extraction of a common scalar is theoretically grounded in the **Prefix Path Property**, which states that for any node a_i in a path P within an FP-tree, the frequency count of every node in its prefix subpath must carry the same count as node a_i [5]. In a quantitative context, this property suggests that items appearing in

a single path often co-occur in fixed ratios, justifying the use of a scaling factor to represent the magnitude of the transaction [4, 5].

Consequently, each pattern P can be uniquely decomposed into its **Canonical Shape** S , which represents the irreducible proportional ratio of the items [4]. The shape S is derived by normalizing each quantity in the pattern by the scaling factor:

$$S = \left\{ \left(i_j, \frac{q_j}{\alpha} \right) \mid (i_j, q_j) \in P \right\} \quad (1)$$

By sorting these items in a fixed lexicographical order, we ensure a unique representation for each structural relationship [4, 6, 7]. Under this formulation, any extracted frequent pattern is represented as a scalar multiple of its base structure: $P = \alpha \cdot S$ [4].

3 Distributed Architecture

The massive scale of modern retail transactions necessitates a distributed computing approach to avoid memory overflows and prohibitive computational costs [3, 4]. Our architecture leverages the **Parallel FP-Growth (PFP)** framework [5], adapted to process quantitative itemsets through a multi-stage MapReduce-based pipeline. This architecture ensures virtually linear speedup by eliminating inter-node communication during the mining process [3, 6].

3.1 Bijective Token Indexing and Vocabulary Discovery

Standard FIM algorithms are designed to process discrete tokens rather than complex tuples. To bridge this gap without modifying the underlying prefix-tree logic, we implement a **Bijective Token Indexing** phase (adapted from PFP’s Parallel Counting step [7, 8]).

- **Sharding:** The raw transactional database DB is partitioned into P independent shards across the cluster [1].
- **Parallel Counting:** Each mapper discovers the local “vocabulary” of $(item, quantity)$ pairs. The MapReduce infrastructure aggregates these counts to generate a global **F-List** [7, 9].
- **Tokenization:** Every unique pair is mapped to a synthetic integer identifier $ID \in \mathbb{Z}^+$. This standardizes the input for the PFP engine, ensuring that a purchase of “2 colas” is treated as a distinct frequent entity from “1 cola” [10].

3.2 Workload Partitioning and the Group List

A critical challenge in distributed FP-Growth is balancing the computational load. PFP addresses this by dividing the $|I|$ items from the F-List into Q independent groups, creating a **Group List (G-list)** [7]. Each group is assigned a unique *group-id* (gid). By setting Q significantly higher than the number of physical

machines (P), we ensure that if one node finishes its mining task early, it can immediately process the next group, preventing “straggler” nodes from bottlenecking the pipeline [11, 12].

3.3 Independent Parallel Mining

The core of the architecture lies in the **Parallel FP-Growth Step** [13]. This stage transforms the database into group-dependent transactions:

1. **Mapper Phase:** Each mapper reads the G-list and outputs group-dependent transactions. For each item in a transaction, it emits a key-value pair where the key is the *gid* and the value is the prefix path of that item [13, 14].
2. **Reducer Phase:** Reducers receive independent “shards” of group-dependent transactions. Each reducer builds a **local FP-tree** and mines it recursively to extract patterns [15, 16].

Because the shards are mathematically independent, there is no need for synchronization between workers during the recursive tree-growth process, which is the primary reason for PFP’s high scalability [3, 5].

3.4 Max-Heap Aggregation

To manage the potentially combinatorial number of patterns, each reducer maintains a **Max-Heap (HP)** of size K [17, 18]. Instead of emitting all discovered patterns, the reducer only keeps the *top-K* mostly supported patterns for each item [19]. This filtered output is then aggregated in the final step to produce the global “Atlas of Proportionality,” which serves as the input for our structural graph reduction [20, 21].

4 Implementation & Algorithm Design

The proposed computational pipeline is engineered to bridge the gap between distributed statistical mining and structural relationship modeling. The implementation is divided into three functional phases, each addressing a specific challenge of quantitative big data processing within an Apache Spark environment.

4.1 Phase 1: Distributed Dataset Transformation (Algorithm 1)

Standard Frequent Itemset Mining (FIM) engines, including the PFP framework, are designed to process discrete, non-quantitative tokens. To preserve purchased quantities without altering the underlying prefix-tree logic, we implement a *Bijective Token Indexing* phase.

This phase executes a distributed transformation where every unique $(item, quantity)$ tuple discovered across the database is assigned a continuous synthetic identifier $ID \in \mathbb{Z}^+$. By leveraging Spark’s

`zipWithIndex` operation, we effectively perform the “vocabulary discovery” described in Step 2 of the PFP framework. This pre-processing ensures that quantitative variations (e.g., “cola:1” vs. “cola:2”) are treated as distinct frequent entities, maintaining the granularity required for proportional analysis.

Algorithm 1: Bijective Token Indexing

Input: Transaction database DB with $(item, quantity)$ tuples
Output: Transformed database DB' with integer IDs
parallel for each transaction $T \in DB$ **extract** all pairs (i, q) ;
aggregate global vocabulary of unique (i, q) pairs;
assign unique $ID \in \mathbb{Z}^+$ using `zipWithIndex`;
parallel for each transaction $T \in DB$ **replace** each (i, q) with corresponding ID ;
return DB' ;

4.2 Local FP-Growth Mining

The local FP-Growth procedure was implemented from scratch rather than delegating to MLlib’s built-in `FPGrowth`, which operates on discrete tokens and cannot handle the quantitative tuples produced by the Bijective Token Indexing phase. The implementation centres on a custom `FPTree` class that maintains an explicit header table — a dictionary mapping each item to the list of its nodes in the tree — enabling direct traversal during conditional tree construction without a full tree scan.

A correctness-critical design choice lies in the conditional pattern base construction: every prefix node is assigned the count of the conditioning item rather than its own count. This ensures that the support of any pattern discovered in the conditional tree is bounded by the frequency of the conditioning item, as required by the Prefix Path Property. Deviating from this — for instance, preserving original node counts — would inflate support estimates and produce spurious frequent patterns.

When a conditional tree reduces to a single path, all nodes inherit the same count from the conditioning item, making support-based ranking among combinations redundant. Our implementation exploits this by enumerating combinations in strictly ascending order of cardinality, retaining only the first K . Beyond correctness, this ordering is a deliberate system-level trade-off: shorter patterns incur lower serialisation overhead during the `combineByKey` shuffle, reducing inter-node communication cost. This is particularly relevant in long-tail datasets where a disproportionate number of conditional trees reduce to single paths.

Algorithm 2: TOPKCOMBINATIONS — cardinality-ascending enumeration

Input: Single-path nodes $\mathcal{P}$, budget K
Output: First K combinations in ascending cardinality order
 $count \leftarrow 0$;
for $r \leftarrow 1$ **to** $|\mathcal{P}|$ **do**
 for each combination C **of size** r **from** $\mathcal{P}$ **do**
 if $count \geq K$ **then return**;
 yield C ;
 $count \leftarrow count + 1$;

Pattern volume is controlled through a Max-Heap of size K , maintained locally by each worker. When the heap exceeds capacity, the lowest-support pattern is evicted. Local heaps are merged in the aggregation phase via `combineByKey`, which ensures the merging is associative and commutative — a requirement for correct distributed reduction. The choice of K directly governs the precision-scalability trade-off: a larger K retains more structural information but increases shuffle volume and memory pressure on the aggregator nodes.

Algorithm 3: Local Top- K FP-Growth

Input: FP-tree T , minimum support ξ , suffix σ , heap size K
Output: Top- K frequent patterns with support
if T **reduces to a single path** $\mathcal{P}$ **then**
 $H \leftarrow []$;
 for each $C \in \text{TOPKCOMBINATIONS}(\mathcal{P}, K)$ **do**
 $sup \leftarrow C[0].count$;
 if $sup \geq \xi$ **then**
 append $(sup, \text{items}(C) + \sigma)$ to H ;
 return H ;
 $\mathcal{H} \leftarrow []$;
for each $(item, nodes)$ **in header table of** T **do**
 $sup \leftarrow \sum_{n \in nodes} n.count$;
 if $sup < \xi$ **then skip**;
 $\pi \leftarrow (item,) + \sigma$;
 $\mathcal{H} \leftarrow \text{ADDTOHEAP}(\mathcal{H}, (sup, \pi), K)$;
 $T' \leftarrow \text{CONDITIONALTREE}(T, item, \xi)$;
 if T' **is non-empty then**
 $\mathcal{H}' \leftarrow \text{TOPKFPKGROWTH}(T', \xi, \pi, K)$;
 $\mathcal{H} \leftarrow \text{MERGEHEAPS}(\mathcal{H}, \mathcal{H}', K)$;
return $\mathcal{H}$;

4.3 Phase 2: Parallel Pattern Extraction (Algorithm 2)

With the database standardized into ID sequences, the system triggers the Parallel FP-Growth (PFP) engine. This stage utilizes the distributed sharding logic where the transaction database is divided into independent parts, allowing each worker to execute mining tasks without inter-machine communication.

To manage the “potential combinatorial number

of candidate patterns” that typically bottlenecks FP-Growth on large datasets, our implementation utilizes a Top- K Max-Heap (HP). As patterns are discovered recursively through conditional FP-trees, they are inserted into a heap of size K indexed by their support value. This ensures that only the most significant patterns are retained, preventing memory overflow. Once mining is complete, an inverse mapping M^{-1} is applied to restore the structural (item, quantity) tuples for final analysis.

Algorithm 4: Parallel FP-Growth with Top- K Filtering

Input: Transformed database DB' , parameter K

Output: Top- K frequent patterns
partition DB' into independent shards;
parallel for each shard build local FP-tree;
mine conditional FP-trees recursively;
maintain Max-Heap HP of size K ;
for each discovered pattern P **do**
 insert into HP based on support;
 prune if $|HP| > K$;

aggregate all local heaps;
apply inverse mapping M^{-1} to recover
(item, quantity) tuples;
return Top- K patterns;

4.4 Phase 3: Proportional Graph Construction (Algorithm 3)

Algorithm 5: Proportional Graph Construction with Transitive Modulo Pruning

Input: Set of frequent patterns $\mathcal{P}$
Output: Directed Acyclic Graph (DAG)
parallel for each pattern $P \in \mathcal{P}$ compute
 $\alpha = \gcd(q_1, \dots, q_k)$;
compute canonical shape $S = P/\alpha$;
group patterns by identical S ;
parallel for each group S let $V = \{\alpha_1, \dots, \alpha_n\}$;
foreach (α_i, α_j) **do**
 if $\alpha_j \bmod \alpha_i = 0$ **then**
 create edge (P_i, P_j) ;
apply transitive reduction: remove (P_i, P_j) if
 $\exists P_k$ such that $P_i \rightarrow P_k \rightarrow P_j$;
return DAG;

The final phase performs the topological clean-up of the mining results, transforming flat lists of frequent itemsets into a hierarchical Directed Acyclic Graph (DAG).

4.4.1 GCD Decomposition

For each extracted pattern, the system calculates the Scaling Factor (α) and the Canonical Shape (S) using a Greatest Common Divisor (GCD) calculation. This approach is theoretically grounded in the Prefix Path Property, which states that nodes in a conditional FP-tree path share a common count inherited from the

conditioning item. By detaching the core ratio (S) from its magnitude (α), we can group patterns that represent the same structural consumer behavior.

4.4.2 Transitive Modulo Pruning

To refine the resulting graph, we apply a custom *Transitive Modulo Pruning* algorithm. Using Spark’s `groupByKey` to cluster matching shapes S onto single worker nodes, the algorithm verifies pairwise integer divisibility among scaling factors:

- **Divisibility Check:** An edge is created if

$$\alpha_j \bmod \alpha_i = 0,$$

signifying a perfect scalar transition.

- **Transitive Reduction:** To eliminate informational redundancy, any direct “shortcut” edge (P_i, P_j) is pruned if an alternative path exists through an intermediate node P_k .

This optimization transforms the raw output into a clean, evolutionary trajectory of purchase behaviors, highlighting direct scaling increments and complex multi-parent convergences while removing redundant links.

5 Performance Analysis

This section evaluates the computational performance and scalability of the proposed pipeline. The primary objective is to demonstrate that the integration of quantitative proportionality analysis does not compromise the high-performance characteristics of the Parallel FP-Growth (PFP) framework.

5.1 Experimental Setup and Dataset Characteristics

To validate the system’s scalability, we reference the performance metrics obtained on Google’s distributed MapReduce infrastructure using the TTD (Tag-Tag Dataset). The TTD represents a massive retail-like environment, containing approximately 15.8 million transactions and over 1 million unique items.

A significant challenge in such datasets is the “long-tail” distribution, where a small number of items appear in millions of transactions, while the vast majority appear with low individual support. Our architecture addresses this by utilizing a low support threshold (ξ) to capture these niche but structural proportional relationships.

5.2 Computational Bottleneck Analysis

The overall execution time is dominated by Step 4: Parallel FP-Growth. In the TTD test, while Step 4 required several thousand seconds, the auxiliary phases

(Sharding, Counting, and Aggregating) required less than 4 seconds combined.

This validates our decision to implement the Proportional Graph Construction (Algorithm 3) as a post-processing step within the Aggregating phase. By confining the complex GCD calculations and the Transitive Modulo Pruning to the aggregated top- K results, the analytical extension avoids adding significant overhead to the recursive mining process, which remains the primary bottleneck.

5.3 Efficiency of the Topological Clean-up

The custom Transitive Modulo Pruning provides a critical “clean-up” that addresses the “combinatorial explosion” of patterns often encountered in retail datasets. Without this pruning, the number of redundant scaling relationships could overwhelm the visualization and analytical layers.

By applying the transitive reduction logic locally on each Canonical Shape group, we transform a flat list of thousands of patterns into a sparse, hierarchical Directed Acyclic Graph (DAG). This optimization ensures that even as the number of extracted frequent patterns increases, the resulting Proportional Relationship Graph remains interpretable and computationally manageable for end-users.

6 Illustrative Application and Case Study

To validate the pipeline’s capability to process non-linear scaling topologies and execute transitive reductions on complex node structures, this section presents a case study based on two distinct consumer behaviors extracted from a simulated retail dataset.

6.1 Scenario Definition: Disconnected Consumer Behaviors

The distributed Parallel FP-Growth engine identified two entirely disconnected frequent structures, representing different purchasing intents:

- **Graph G_{S_1} (Beverage Party Pack):** A multi-item association between cola, juice, and water with a base ratio of 1 : 2 : 3.
- **Graph G_{S_2} (Pancake Breakfast Bundle):** A two-item association between flour and milk with a base ratio of 2 : 1.

The extraction of these two distinct shapes demonstrates the system’s ability to maintain structural isolation through Spark’s partitioning logic, ensuring that unrelated market baskets do not interfere with each other’s topological reduction.

6.2 Structural Decomposition and Scale Extraction

The pipeline executes a localized parallel Greatest Common Divisor (GCD) calculation for each pattern. Table ?? summarizes the decomposition of these patterns into their Canonical Shapes (S) and Scaling Factors (α).

This decomposition confirms the Prefix Path Property: although the absolute quantities vary, the core structural ratio remains constant, allowing the system to group these patterns into single evolutionary trajectories.

6.3 Multi-Parent Convergence and Topological Pruning

The Beverage cluster (G_{S_1}) serves as a rigorous test for the Transitive Modulo Pruning algorithm. The scale set $V_{S_1} = \{1, 2, 3, 6\}$ presents a complex divisibility structure:

- The pattern P_4 ($\alpha = 6$) is a perfect scalar multiple of all previous nodes (P_1, P_2, P_3).
- The raw edge generation would yield redundant shortcuts, such as the direct link (P_1, P_4).

Our algorithm identifies that P_4 is reachable from P_1 through intermediate steps ($1 \rightarrow 2 \rightarrow 6$ or $1 \rightarrow 3 \rightarrow 6$). Consequently, the redundant edge (P_1, P_4) is pruned, while both (P_2, P_4) and (P_3, P_4) are preserved. This is because $2 \nmid 3$ and $3 \nmid 2$, representing independent “evolutionary paths” toward the larger pack.

6.4 Analytical Insights for Retail Intelligence

The resulting Proportional Relationship Graph provides two key insights:

- **Linear Growth Identification:** G_{S_2} shows a simple doubling behavior ($P_5 \rightarrow P_6 \rightarrow P_7$), suggesting a predictable scaling in breakfast bundle purchases.
- **Structural Convergence:** G_{S_1} reveals that the large-scale “Party Pack” (P_4) is not just a bulk version of the base unit, but a point of convergence for different purchasing behaviors.

For a retailer, this optimized Directed Acyclic Graph (DAG) identifies the most frequent “natural increments” in purchase volume, enabling more precise bundle recommendations and targeted upselling strategies compared to traditional, non-proportional FIM outputs.

7 Limitations and Conclusion

7.1 Limitations and Future Work

While the proposed pipeline successfully extracts structural proportions from massive datasets, several con-

Graph Partition	Pattern	$g = \gcd(\cdot)$	Canonical Shape (S)	Scale (α)
G_{S_1} (Beverage)	P_1	$\gcd(1, 2, 3) = 1$	$\{(\text{cola}, 1), (\text{juice}, 2), (\text{water}, 3)\}$	1
	P_2	$\gcd(2, 4, 6) = 2$	$\{(\text{cola}, 1), (\text{juice}, 2), (\text{water}, 3)\}$	2
	P_3	$\gcd(3, 6, 9) = 3$	$\{(\text{cola}, 1), (\text{juice}, 2), (\text{water}, 3)\}$	3
	P_4	$\gcd(6, 12, 18) = 6$	$\{(\text{cola}, 1), (\text{juice}, 2), (\text{water}, 3)\}$	6
G_{S_2} (Pancake)	P_5	$\gcd(2, 1) = 1$	$\{(\text{flour}, 2), (\text{milk}, 1)\}$	1
	P_6	$\gcd(4, 2) = 2$	$\{(\text{flour}, 2), (\text{milk}, 1)\}$	2
	P_7	$\gcd(8, 4) = 4$	$\{(\text{flour}, 2), (\text{milk}, 1)\}$	4

Table 1: GCD-based decomposition into canonical shapes and scaling factors




Figure 1: Comparison of proportional relationship graphs: (a) multi-parent convergence and (b) linear scaling behavior

straints remain that offer opportunities for future research:

- **Sensitivity to Quantity Noise:** The reliance on the Greatest Common Divisor (GCD) for extracting the Scaling Factor (α) makes the model highly sensitive to “noise” in purchase quantities. For instance, a minor deviation in a single transaction (e.g., a user purchasing 11 units instead of 10) can prevent the discovery of a 1:2 ratio, as the integer factorization logic is strict. Future iterations could incorporate a “fuzzy” GCD or clustering-based approach to handle near-proportionality.
- **Strict Proportionality Constraints:** The current model assumes that consumer behavior scales linearly ($P = \alpha \cdot S$). It does not account for non-linear trends, such as “bulk-buy” behaviors where product ratios might shift as total volume increases.
- **Absence of Temporal Analysis:** The pipeline processes the dataset as a static snapshot. It does not currently capture seasonal variations or time-based shifts in purchase proportions, which are critical for dynamic retail environments.

7.2 Conclusion

This work has successfully extended the traditional Frequent Itemset Mining (FIM) paradigm by introducing a distributed framework for proportional quantitative analysis. By leveraging the Parallel FP-Growth (PFP) engine, we have demonstrated that it is possible to move beyond simple binary associations to uncover

the underlying structural ratios of consumer behavior without sacrificing computational scalability.

Our implementation introduced three key technical contributions:

- **Bijective Token Indexing:** A transformation layer that allows the PFP distributed engine to process quantitative tuples while maintaining compatibility with standard prefix-tree structures.
- **Structural Decomposition:** A mathematical approach, anchored in the Prefix Path Property of FP-trees, that utilizes integer factorization to detach a pattern’s irreducible Canonical Shape from its scalar magnitude.
- **Topological Synthesis:** A post-processing phase that applies Transitive Modulo Pruning to eliminate redundant informational pathways, transforming a flat list of patterns into an optimized Directed Acyclic Graph (DAG).

The performance analysis, supported by empirical data from large-scale web datasets, confirms that this architecture achieves virtually linear speedup, making it suitable for industrial-scale retail applications. By providing an “Atlas of Proportionality” rather than a simple list of popular items, this system empowers businesses with deeper insights for automated bundle recommendations and targeted inventory management, effectively bridging the gap between big data mining and actionable business intelligence.